Predictive Risk and Recording Decisions in Partially Observed Neural Networks

Jiayi Du[1*†] and Shuaizhi Cheng[2†]

[1*]Stiftung Louisenlund, Germany.
[2]Department of Computing, Harbin Institute of Technology, China.

*Corresponding author(s). E-mail(s): jiayi.du@mail.louisenlund.de;
†These authors contributed equally to this work.

Abstract

Limited recordings constrain how well a fitted neural network predicts a planned stimulus response. We ask which calibration information supports that prediction, and whether averaging a local predictive-risk criterion over physiological uncertainty improves recording decisions.

In a 100-unit rank-two model with known connectivity and uncertain neuronal gains, allocating five recordings to the stimulation site, three readouts and one additional site reduced response MSE by 96.44% relative to five nonendpoint recordings in 40 networks. An independent 40-network validation gave a change of −98.13% (95% CI −99.19% to −93.94%). All recordings concerned calibration activity; the predicted stimulus response was withheld.

A further 40-network study examined the mechanism using alternative parameter fits and controlled information changes. Joint endpoint recording changed the mean squared width of feasible response ranges by -99.08% (95% CI -99.62% to -96.10%) and response MSE by -96.22% (95% CI -99.63% to -79.45%). Degrading endpoint recording precision broadened response ambiguity. Revealing the four endpoint gains and refitting the others changed response MSE by -98.03% (95% CI -99.26% to -92.77%) relative to revealing four ordinary gains. These results support a task-specific role for endpoint information in constraining the predicted response, while individual parameters can remain uncertain.

The original 12-unit comparisons and independent 100-unit confirmation left the additional benefit of prior averaging uncertain. Two covariance-based recording criteria were also evaluated on 40 disjoint networks with endpoints inaccessible. The internal covariance criterion changed MSE by -22.77% (95% CI -52.62% to -5.70%) relative to nominal selection in a secondary comparison, but this advantage weakened under additional fitting starts. The prespecified comparison did not establish an improvement from the held-out-prior criterion. The contribution is a replicated recording-allocation finding and a controlled account of its information mechanism within the tested model family.

1 Introduction

A recording set should be judged by how well its data support the intended prediction. Biohybrid robots translate optical inputs into motion through directly stimulated muscle or an intervening neuromuscular interface [1, 2] (reviewed in [3]), and earlier neuromuscular bots coupled cultured neurons to muscle tissue [4]; cultured neural networks also support decoding of responses to patterned electrical inputs [5], and closed-loop use of cultured networks through a small electrode set has a longer history [6]. These systems motivate a task-specific view of neural measurements, although a neural response is not itself a mechanical output. The inference may concern population activity or the response of a readout to a planned stimulus, and under a limited recording budget these tasks value the same observation differently. A recording-design evaluation must therefore connect the design criterion to the prediction made from the selected data. This paper studies that connection in simulation only: the networks are generic rate models, not a model of any biohybrid platform.

Connectivity constrains neural dynamics but does not fix all neuronal properties. Different circuit parameters can produce similar activity [7], and systems-biology models show that uncertain parameters can coexist with well-constrained predictions [8]. Beiran and Litwin-Kumar study connectome-constrained activity prediction with uncertain physiological parameters, partial recordings and recording-site selection [9]. Lappalainen and colleagues combine anatomical structure, simplified dynamics and task optimisation to predict visual responses [10]. These studies motivate asking how recording decisions support a specified held-out intervention response. Here we examine both the incremental value of prior averaging and the information supplied by recording the stimulation and readout sites during calibration.

Validation must target the held-out response. Channel selection in brain-computer interfaces is usually posed as choosing channels that maximise decoding performance on recorded data [11]; here the choice is made before recording, from a model. Even high-density probes expose more sites than can be read out at once [12], so a choice of recorded sites is still required; biohybrid preparations may be more restricted. Sensor-selection studies in nonlinear networks distinguish structural observability from the practical accuracy of state estimation [13], and budgeted sensor placement is commonly solved by greedy selection with guarantees for suitable objectives [14]. These works treat the selection problem well, but the guarantees concern the selection objective, not the error of a response predicted after fitting. Agreement with the activity used for fitting does not by itself establish accuracy for an unobserved perturbation. Sourmpis and colleagues use perturbation outcomes to assess neural models and the contribution of biological constraints [15]. Here we fix the stimulation protocol and ask which calibration recordings support its response prediction. The variable being designed is the recording set, rather than the stimulation input, which has its own experimental-design literature [16, 17].

Bayesian experimental design values observations by the expected utility of the resulting inference [18, 19], and goal-oriented variants value them according to a quantity of interest [20]. A nominal-parameter approximation is inexpensive; averaging local criteria over a parameter ensemble, as in related designs, accounts for variation in sensitivities [21], while more complete nonlinear approaches address a broader inference problem [22–25]. These designs refine the risk, but local sensitivity-based risk need not track the error of a nonlinear response predicted after fitting. Two benefits must be distinguished: forecasting the prediction error more accurately than the nominal risk, and changing the selected observations enough to reduce that error. A revised risk estimate can leave the recording set unchanged. Conversely, a design specialised for one known readout can outperform a shared design through task specialisation alone. The study therefore compares criteria that differ in only one of these two choices. Hypothesis: prior averaging improves response prediction or the forecast accuracy of the risk; both routes require that the prior-averaged criterion differ from the nominal criterion in the recordings it selects or in the forecast it reports. Under a shared fitting procedure, changing the recording set can change the point prediction; changing the reported risk alone cannot. Decision effects and forecast accuracy can therefore be evaluated separately.

A controlled computational study evaluates this question. Scalar and vector response objectives are crossed with nominal and prior-ensemble local-risk criteria, using the same model information, feasible pool, recording budget and estimator. We first compare recording decisions and actual response error, then evaluate risk forecasts on fixed predictions. An independent 100-unit confirmation tests the pilot result in a larger model family. Focused follow-ups vary uncertainty within that family and replace random recording slots with stimulation and readout sites. An additional independent cohort tests whether endpoint calibration narrows response disagreement among comparably fitted models, whether

that disagreement increases as endpoint recording precision falls, and whether revealing endpoint gains is more informative than revealing equally many ordinary gains. A separate indirect-recording comparison evaluates two covariance-based criteria. These follow-ups connect the original criterion comparison to the information that improves prediction and its mechanism.

The study contributes a controlled comparison separating the prediction objective from the risk approximation; a replicated, equal-budget recording-allocation result identifying the value of covering both stimulation and readout sites; and an independent investigation of how that information constrains the predicted response. Response-range searches, recording-precision interventions and gain-information controls connect the allocation benefit to predictive ambiguity. Established local-risk, covariance and MAP procedures provide the methodological tools. Their task-specific evaluation and the controlled information mechanism are the empirical contributions.

2 Results

2.1 A common information budget separates the prediction objective from the risk approximation

The simulated networks have known numerical connectivity J and uncertain neuronal gains and biases (symbols in the Notation table, Methods). Calibration data were generated from a limited set of nodes; the evaluation target was the time-integrated response to a separate fixed stimulus, relative to an unstimulated control (Fig. 2). Three readouts were chosen using structural reachability before generating physiological parameters or measurement noise. The scalar objective specialised its recording set for one readout, whereas the vector objective used one set for the same three-response family. This construction separates task specialisation from the comparison between nominal and ensemble risk. Throughout the Results, "single-response" (scalar) and "shared" (three-response, vector) name the two objectives; the abbreviations LS, ES, LV and EV, defined in Methods, combine objective and risk criterion.

Box 1 and Fig. 1 summarise the audit. All methods fitted the selected observations with the same Gaussian MAP estimator and predicted the response by nonlinear simulation. The principal candidate pool excluded the stimulation site and all three readouts, preventing direct endpoint recording. All-site recording and a larger budget were retained as additional conditions. Exact dataset reuse avoided redundant fits when methods selected identical nodes; no observations from different scalar designs were pooled. Box 1 (below) defines the three-step audit that structures the Results: Decision, Error and Forecast.

Box 1 | Three-step audit of a recording criterion. Inputs: a design pool, a budget M, a stimulus and target q, an estimator, and two criteria A (nominal) and B (prior-averaged). 1. *Decision.* Select a recording set with each criterion. If the sets are identical the fitted data, predictions and errors are identical, so only the remaining designs can differ. 2. *Error.* Fit with the common estimator and compare response error against hidden truth across all designs, including those with identical recording sets, paired by independent network, with a cluster-bootstrap interval. 3. *Forecast.* Fix the recording set, data and point prediction; compare the two risk forecasts of the error with a score that is minimised at the true mean squared error, again paired by network.

Report agreement counts, intervals read as approximate bounds, and the influence of the largest single network.

Before the methods were compared, the code was checked against an independent implementation and a linear-Gaussian case with a known answer. An independent rate/sensitivity implementation agreed within approximately 3.4×10^{-16}, and finite differences supported the analytic derivatives. A nondegenerate linear-Gaussian positive control gave identical nominal and ensemble designs in all 48 design pairs; analytic prediction risk agreed with Monte Carlo MAP error across 144 target/design evaluations (Supplementary Fig. 5). A separate additive-bias control produced the expected zero derivative of the paired stimulus response. These checks validate specific computational operations, rather than biological assumptions.

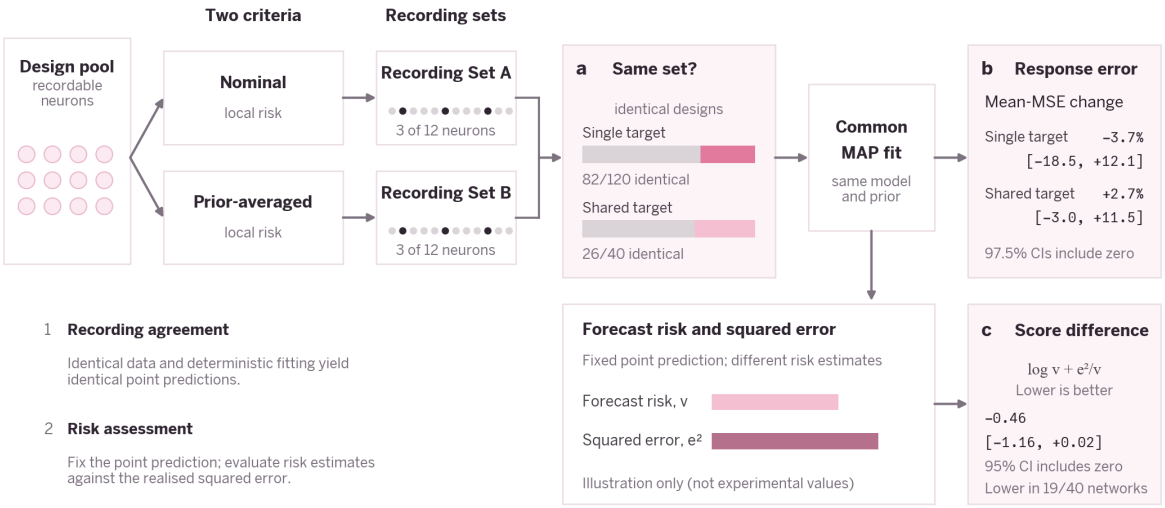


Fig. 1 | The three-step audit. Two criteria, nominal and prior-averaged local risk, each select a recording set from the same design pool. Panel a asks whether the two sets are identical; if they are, the common maximum a posteriori fit sees the same data and the response error cannot differ. Panel b compares the response error of the fitted prediction against hidden truth (mean-MSE change with the 97.5% interval for the single-response and shared objectives; both intervals include zero). Panel c fixes the data and the point prediction and compares the two risk forecasts v of the error e with the score $\log v + e^2/v$. The reported mean score difference is prior-averaged minus nominal (95% interval includes zero). The counts, effects and intervals are those of Figs. 3–4 and Table 1; the dot and bar icons for the recording sets and for v and e^2 are illustrative and carry no data. Negative differences favour the prior-averaged criterion.

2.2 Prior averaging often selects the same recordings and did not establish lower error in the primary comparisons

The Decision and Error steps of Box 1 apply to the first frozen evaluation, which used 40 independent J matrices, with three physiological truths, three noise repetitions and three readouts nested within each network. The primary condition combined broad matched uncertainty, excluded endpoints and a budget of three recorded nodes. For each method, errors were first averaged within J; mean risks and paired intervals were then calculated across networks. All fits terminated successfully, with no fallback or invalid task scale.

The two criteria often chose the same recordings (Fig. 3a). A descriptive readback of the frozen designs found identical recording sets in 82 of 120 target-specific designs (68%; three designs per network) and in 26 of 40 shared designs (65%). Because identical sets give identical fitted data and predictions by construction, a difference in error can arise only in the remaining designs. At the network level, the ensemble and nominal MSE were exactly equal in 39 of 80 network–objective cells (13 single-response and 26 three-response networks; Fig. 3b). This count is descriptive and was not a prespecified hypothesis.

The ensemble scalar design had a slightly lower mean MSE than its nominal counterpart, but the paired interval included both improvement and deterioration (Table 1; Fig. 3c). The vector comparison had the opposite point estimate and was also uncertain. Relative mean-MSE changes were −3.74% for ES versus LS and +2.69% for EV versus LV. A leave-one-network diagnostic changed the sign of the scalar mean difference when one influential network was omitted; that network remains in every primary estimate. The results do not establish a stable ensemble advantage or method equivalence.

Recording constraints changed the achievable error. For nominal scalar design at M=3, RMSE was 0.000929 with all sites available and 0.004046 with endpoints excluded under matched uncertainty; the corresponding broad-condition values were 0.001946 and 0.008312. These exploratory comparisons make direct endpoint access an important part of the task definition. The activity and random references, both budgets, and all uncertainty conditions are retained in Supplementary Figs. 2–4.

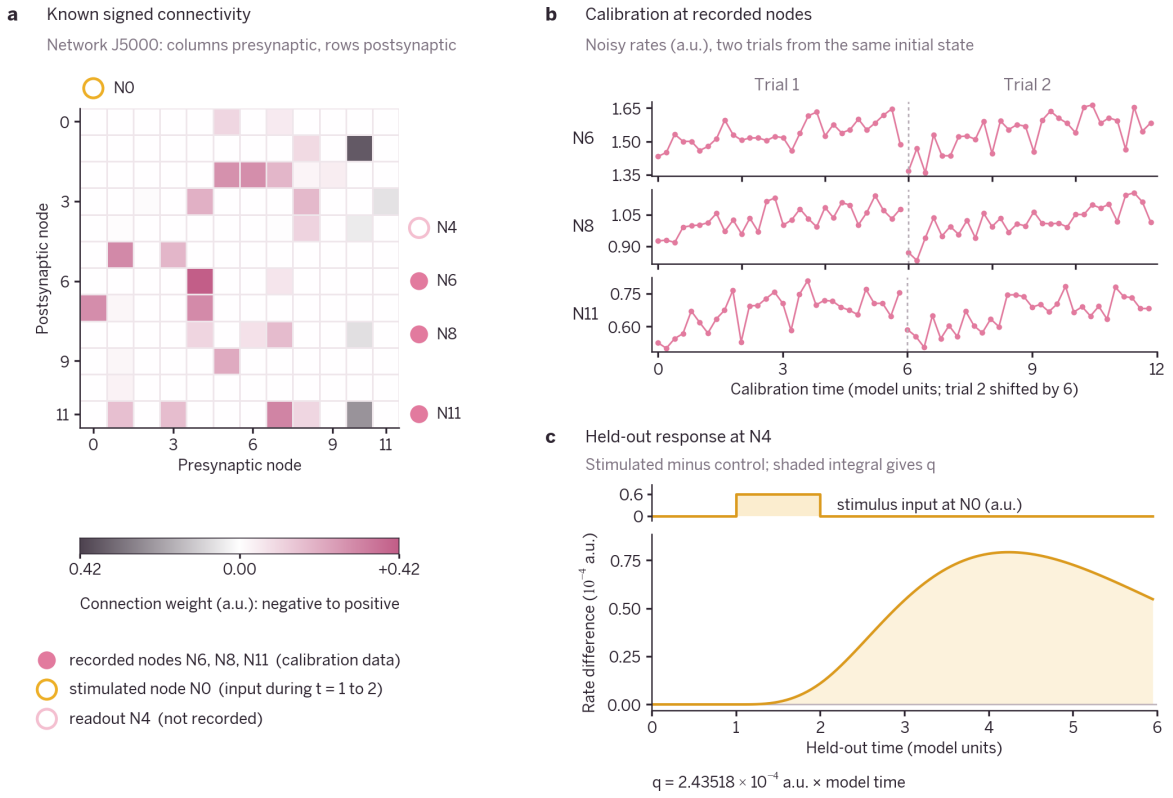


Fig. 2 | Connectivity, limited recordings and a held-out response in one simulated network. a, Known signed connectivity J for the first frozen network, J5000 (rose positive, deep plum negative weights; symmetric colour limits ±0.42 around exact zero); columns are presynaptic and rows postsynaptic nodes. Filled markers beside the matrix identify the recorded nodes N6, N8 and N11, the open light-rose marker the unrecorded readout N4, and the open amber marker the stimulated node N0. **b,** Actual noisy calibration observations at the three recorded nodes in two trials, each reset to the same known initial state (Gaussian observation noise SD 0.05, stride 4). Trial 2 is shifted by 6 model-time units for display; the dashed line marks the trial boundary, not continuous evolution, and lines connect observations only within a trial. Vertical scales differ between nodes. **c,** Stimulated minus matched zero-input control rate at N4 (displayed in units of 10^{-4}); the upper strip shows the stimulus input at N0 during model time 1 to 2. The integrated response $q = 2.43518 \times 10^{-4}$ (a.u. × model time) is Δt times the sum of the 120 frozen-model differences ($\Delta t = 0.05$), not a new trapezoidal integral. The case is the fixed first-index instance (broad uncertainty, truth and noise draw 0, target 0, endpoints excluded, $M = 3$, nominal single-response selection), chosen before its effect was inspected. Independent $n = 1$ network; no interval is shown and no between-method effect or biological validity is claimed. All quantities are in dimensionless model units.

Table 1 | Frozen primary response-error comparisons

Objective	Nominal RMSE	Ensemble RMSE	Mean MSE difference, E−N	97.5% interval	Relative mean-MSE change [97.5% interval]
Single response	0.008312	0.008155	-2.586×10^{-6}	$[-1.471\times10^{-5}, 4.826\times10^{-6}]$	−3.74% [−18.48%, 12.12%]
Three responses	0.008451	0.008564	1.924×10^{-6}	$[-1.148\times10^{-6}, 6.826\times10^{-6}]$	2.69% [−2.95%, 11.51%]

Broad matched uncertainty; endpoints excluded; M=3; 40 independent networks. RMSE is the square root of the overall mean network MSE. The primary estimand is the raw paired mean-MSE difference. The relative effect is a ratio of aggregate means minus one, not an average of individual percentage changes. Each prespecified comparison uses a 97.5% paired network-cluster interval.

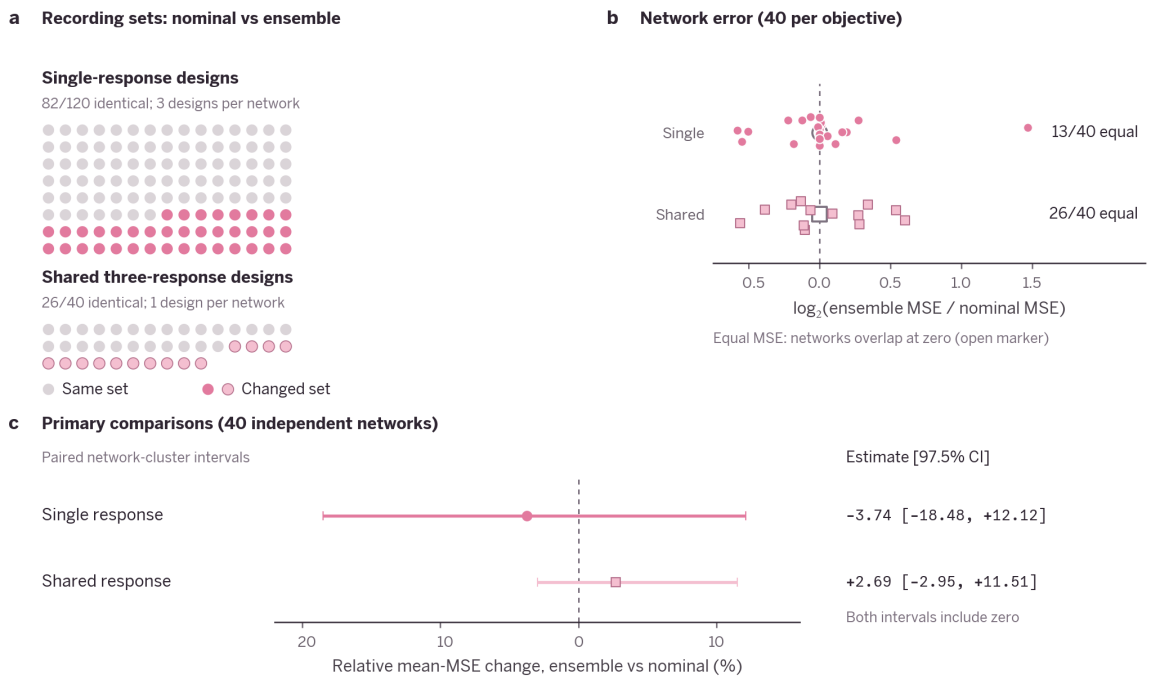


Fig. 3 | Prior averaging often selects the same recordings; primary error comparisons remain uncertain. a, Each dot is one recording design from the frozen first evaluation (120 single-response designs, three per network; 40 shared designs). Grey dots are designs in which the nominal and ensemble criteria chose identical recording sets; coloured dots are designs in which they differed (82/120 and 26/40 identical). **b,** Descriptive within-network $\log_2$ ratio of ensemble to nominal MSE for the 40 independent networks per objective; the hollow marker at zero stands for networks with exactly equal MSE (13 single-response, 26 three-response), and filled points are the remaining networks. Three truths, three noise repetitions and three readouts are averaged within each network. These ratios are neither the primary estimand nor the ratio of aggregate means. **c,** Relative change in aggregate mean MSE, ensemble versus nominal, with the prespecified 97.5% paired network-cluster bootstrap intervals (same estimates as Table 1). Negative values favour the ensemble. All networks, including influential instances, are retained; Supplementary Fig. 1 shows the complete raw difference distributions. No method-equivalence margin was specified.

2.3 A separately specified forecast test separates risk estimation from point-prediction accuracy

In the first study, an exploratory aggregate comparison suggested that ensemble risk could match actual error better than nominal risk does, even when it did not improve the selected recordings. At the primary nominal scalar recording set, actual MSE divided by nominal forecast risk was approximately 1.615; the ratio using ensemble forecast risk was 1.103. Such a ratio does not establish forecast accuracy across networks. The Forecast step of Box 1 therefore became a separate test, specified by us, using 40 new J matrices, fixing the recording policy, selected data and fitted point predictions while comparing the two risk forecasts.

The primary ensemble-minus-nominal MSE-forecast score difference was −0.45565 (95% paired-cluster interval, −1.15857 to 0.02367; Fig. 4). Nineteen network means favoured the ensemble forecast and 21 favoured the nominal forecast. The shared point predictions had MSE 1.1781×10^{-4}, so no point-prediction improvement can be attributed to changing the forecast alone. The interval does not establish a stable forecast-score gain. The mean is also dominated by one network: its score difference of −11.67 accounts for about 64% of the summed differences, and the mean of the other 39 is −0.17. In aggregate, the nominal forecast risk (1.19×10^{-4}) was closer to the actual MSE (1.18×10^{-4}) than the ensemble forecast risk (1.92×10^{-4}), so this study does not show that the ensemble forecast is better calibrated than the nominal forecast. The exploratory aggregate hint from the first study (ensemble ratio 1.10 against nominal 1.62) therefore did not replicate. No fit failures, fallbacks or forecast-floor clipping occurred.

A secondary hidden-bias-shift comparison favoured the ensemble forecast by −1.3583 score units (pointwise 95% interval, −2.2143 to −0.6345). Both mean risk forecasts nevertheless remained below

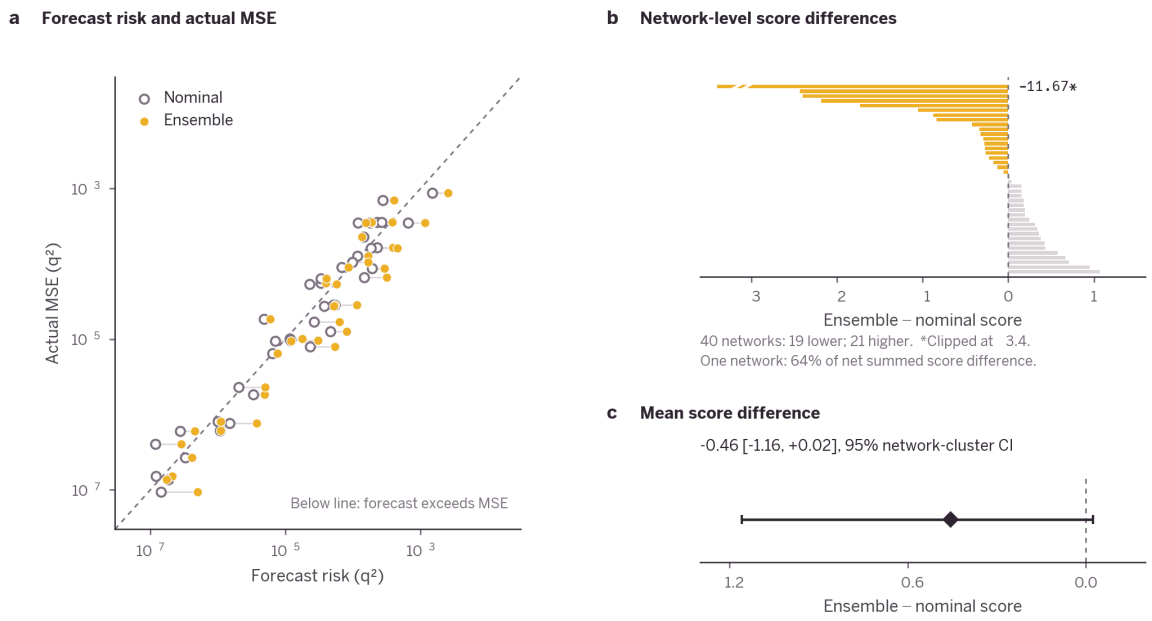


Fig. 4 | Risk forecasts are evaluated on identical predictions. The independent follow-up uses the nominal scalar recording policy, broad matched uncertainty, excluded endpoints and M=3 in 40 new networks. **a,** Nominal (open) and ensemble (filled) mean forecast risks are paired with the same actual mean response MSE for each network; horizontal segments link the two forecasts. Axes have identical logarithmic ranges and a dashed identity line; points below the line have a forecast that exceeds the actual error. This network-aggregate comparison is descriptive, not a stand-alone calibration test. **b,** All 40 network-level score differences, sorted; 19 favour the ensemble forecast (amber) and 21 the nominal forecast (grey). One bar (−11.7) is clipped at the axis limit and labelled. **c,** The mean score difference with the prespecified 95% paired network-cluster interval. Panels b and c use different horizontal scales because individual differences and uncertainty about their mean are distinct statistical objects. Scores $S(v,e)=\log(v)+e^2/v$ are computed for individual target/truth/noise cases and then averaged within network; they are not scores of the aggregate coordinates in a. Negative differences favour the ensemble. The score evaluates an MSE functional, not a complete predictive distribution.

actual MSE in that condition. This result describes a relative forecast improvement in one secondary setting; it does not show that the hidden shift was learned or that uncertainty was fully calibrated.

2.4 Independent confirmation narrows the benefit suggested by the 100-unit pilot

An exploratory extension used eight independent 100-unit rank-two tanh networks from a published mathematical family [9], with adapted calibration inputs, measurement noise and inference. It suggested lower response error from prior averaging in several settings, including a −13.56% shared-objective change at five recordings with endpoints excluded (95% CI −28.56% to −4.78%). We therefore fixed that contrast as the primary comparison in 40 new networks before inspecting their outcomes. The original pilot and new confirmation remain separate sampling stages.

In the confirmation, prior averaging changed shared-objective mean MSE by +7.95% (95% CI −23.59% to +20.95%; 17 of 40 networks favoured the ensemble). Thus, the pilot improvement was not established in the larger independent sample. At ten recordings, large fitting errors made the aggregate contrast especially sensitive to individual networks. The complete objective, budget and pool results are retained in the supplement.

2.5 Equal-budget endpoint calibration yields a reproducible positive result

When all sites were available and five were recorded, nominal and ensemble shared-objective selectors reduced MSE relative to random recording by 95.95% and 88.17%, respectively. Inspection of the selected sets suggested a concrete source of this benefit: both selectors always recorded the stimulation site, and they included all three readouts in 30 and 38 of the 40 networks. We therefore introduced a simple endpoint baseline and controlled which endpoints occupied the same five recording slots.

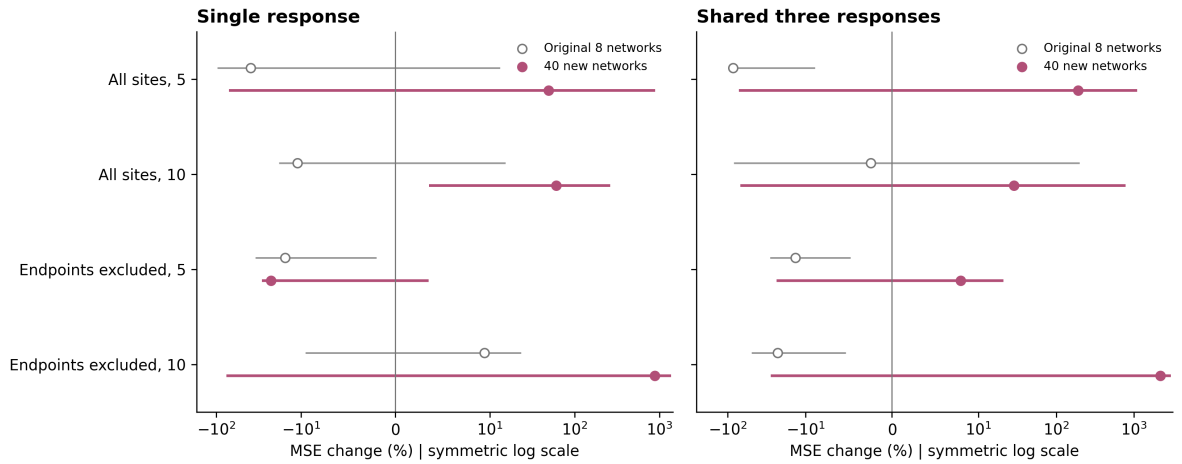


Fig. 5 | Pilot and independent confirmation in the 100-unit family. The eight-network pilot and the 40-network confirmation are plotted separately. Relative changes compare prior-averaged with nominal selection, paired within network. The prespecified confirmation contrast is the shared objective with endpoints excluded and five recordings; other conditions are exploratory. Changing from the 12-unit model to this family also changes dynamics and parameterisation.

Four allocations recorded neither endpoint group, the stimulation site only, the three readouts only, or both groups. Unfilled slots were sampled from the remaining neurons using shared random permutations, with three allocation repeats. All four allocations used the same hidden networks, calibration data, recording count and original estimator. In the first 40 networks, recording both groups changed MSE by -96.44% (95% CI -98.38% to -92.52%) relative to recording neither, with improvement in 38 networks. Relative to stimulation-only and readouts-only allocations, the changes were -95.41% (95% CI -97.79% to -91.69%) and -95.10% (95% CI -97.73% to -90.10%).

A prospective local protocol then fixed the both-versus-neither contrast and evaluated 40 additional, disjoint networks. The primary change was -98.13% (95% CI -99.19% to -93.94%), with improvement in 37 of 40 networks. The secondary changes relative to stimulation-only and readouts-only allocations were -99.28% (95% CI -99.71% to -93.01%) and -95.48% (95% CI -97.58% to -90.42%). Every allocation records calibration activity; the stimulus response being predicted is withheld from fitting.

The simple endpoint baseline also changed the interpretation of the risk-based selectors. In the first 40 networks, nominal selection versus this baseline gave +23.64% (95% CI -41.31% to +161.94%), and ensemble selection gave +260.80% (95% CI -56.39% to +1250.22%). These intervals leave their incremental benefit over endpoint recording uncertain. The positive result is the recording allocation itself: covering both ends produces a large improvement at fixed budget, and that improvement replicates on new networks.

2.6 Uncertainty changes decisions while fitting remains a separate source of variation

Within the same 40 confirmation networks, prior and truth standard deviations were varied together over 0.3, 0.6 and 0.9, with common standardised draws. The mean correlation between nominal and ensemble candidate rankings at the first selection step fell from 0.944 to 0.687 to 0.600. All 40 networks showed a decrease between the narrowest and widest distributions. At five recordings, the average fraction of shared selected sites fell from 0.735 to 0.580 to 0.555. Thus uncertainty substantially changes recording choices in this family, while the corresponding response-error comparisons remain variable. Component controls separating target sensitivities and observation information did not establish one consistently dominant source of error reduction.

We also revisited the primary five-recording comparison using one common fitting rule: the original start, two nearby starts and two prior draws, retaining the smallest MAP objective. The selection rule used calibration data and the prior, with no access to response truth. Across the same 40 networks, ensemble versus nominal selection changed MSE by -49.40% (95% CI -70.90% to +3.65%). For nominal

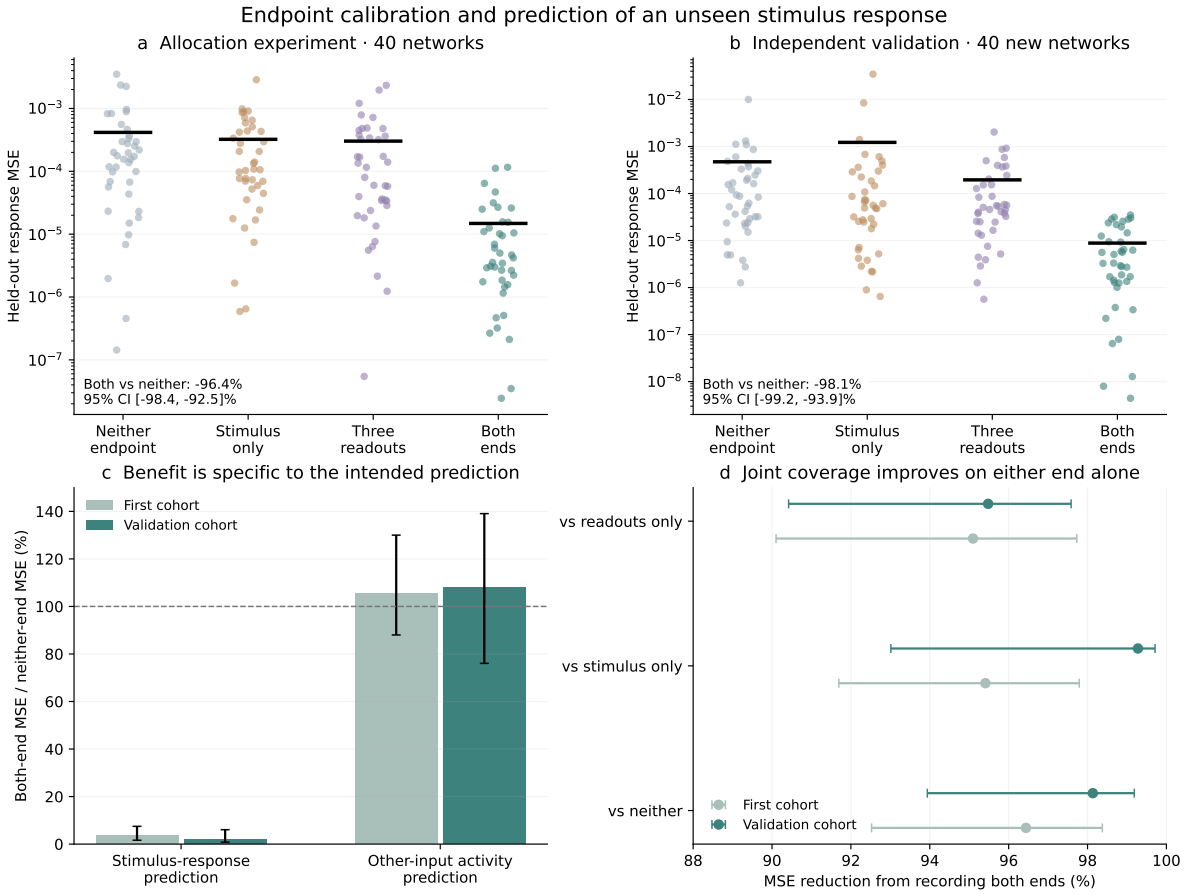


Fig. 6 | Endpoint allocation and held-out response prediction. All four allocations record five sites. Panels a and b show network mean response MSE in the first 40 networks and 40 independent validation networks; dots are individual networks and horizontal bars are arithmetic means. Panel c compares the both-end allocation with the neither-end allocation for the intended stimulus response and for whole-network activity under other held-out inputs; bars show ratios of mean MSE and 95% paired network-bootstrap intervals. The other-input analysis is exploratory. Panel d shows MSE reductions and intervals for joint coverage relative to each partial allocation, with a magnified reduction axis. Both-versus-neither is the primary validation comparison; the other validation contrasts are secondary.

and ensemble selection separately, changes from the original fit were -3.07% (95% CI -13.14% to -0.27%) and -54.56% (95% CI -76.12% to +2.26%). One network contributed 80.7% of the absolute between-selector difference after refitting; three nominal fits still had objectives above the objective at the generating parameters. All 320 retained best fits terminated successfully. These are retrospective estimator-sensitivity results, reported alongside the frozen confirmation rather than replacing it.

Endpoint recording improved functional predictions much more clearly than it recovered individual gains. In the first cohort, mean stimulation-gain MSE changed from 0.661 to 0.579 and mean readout-gain MSE from 0.787 to 0.760 between the neither and both allocations. The validation cohort gave corresponding changes from 0.700 to 0.287 and from 0.743 to 0.507. A nominal local calculation predicted marginal variance reductions from approximately 91% of prior variance to 1–2%, which should be distinguished from these nonlinear parameter-estimation errors. In a separate oracle diagnostic on the first cohort, replacing fitted stimulation and readout gains by their true values, while leaving all other fitted parameters fixed, changed the neither-allocation MSE by -45.45% (95% CI -91.95% to +36.98%). These initial diagnostics motivate the next experiment: endpoint observations can change inferred recurrent dynamics and parameter compensation, so the relevant question is how far response predictions can vary among comparably fitted models.

The improvement was specific to the intended response. Under different held-out calibration-style inputs, whole-network activity MSE changed by +5.36% (95% CI -12.00% to +30.00%) in the first

cohort and +8.05% (95% CI -23.95% to +39.06%) in the validation cohort, despite the large response-error reductions. In an exact descriptive decomposition of response error into a readout-gain term and a dynamic-response term, their negative cross term offset 93.8% and 83.7% of the sum of their separate squared terms under joint endpoint recording, compared with 23.5% and 2.1% under neither-end recording. This supports studying the predictive combinations constrained by calibration, rather than interpreting parameter accuracy alone as the mechanism.

2.7 Endpoint information constrains response ambiguity across comparably fitted models

Forty additional networks, disjoint from both endpoint cohorts, tested whether endpoint observations narrow the responses permitted by comparably fitted parameter vectors. Every allocation used five recordings and a common five-start MAP estimator. For each readout we searched for the smallest and largest nonlinear response within a MAP-objective allowance of 2 above the best fit. These numerical feasible spans describe prediction ambiguity under a common data-and-prior constraint; they are not calibrated confidence intervals.

Relative to five nonendpoint recordings, joint endpoint recording changed the mean squared response-range width by -99.08% (95% CI -99.62% to -96.10%), with contraction in 40/40 networks. Its response-MSE change was -96.22% (95% CI -99.63% to -79.45%), with improvement in 38/40 networks. Raising the standard deviation of endpoint recording noise to four and sixteen times its original value changed the range metric by +471.43% (95% CI +178.12% to +999.62%) and +3384.93% (95% CI +1215.52% to +8268.23%), respectively, relative to ordinary-precision endpoint recording. The corresponding MSE changes were +156.17% (95% CI +24.32% to +928.87%) and +570.70% (95% CI +125.13% to +3169.11%). Noise levels were included correctly in the likelihood, and standardised noise draws were paired across conditions.

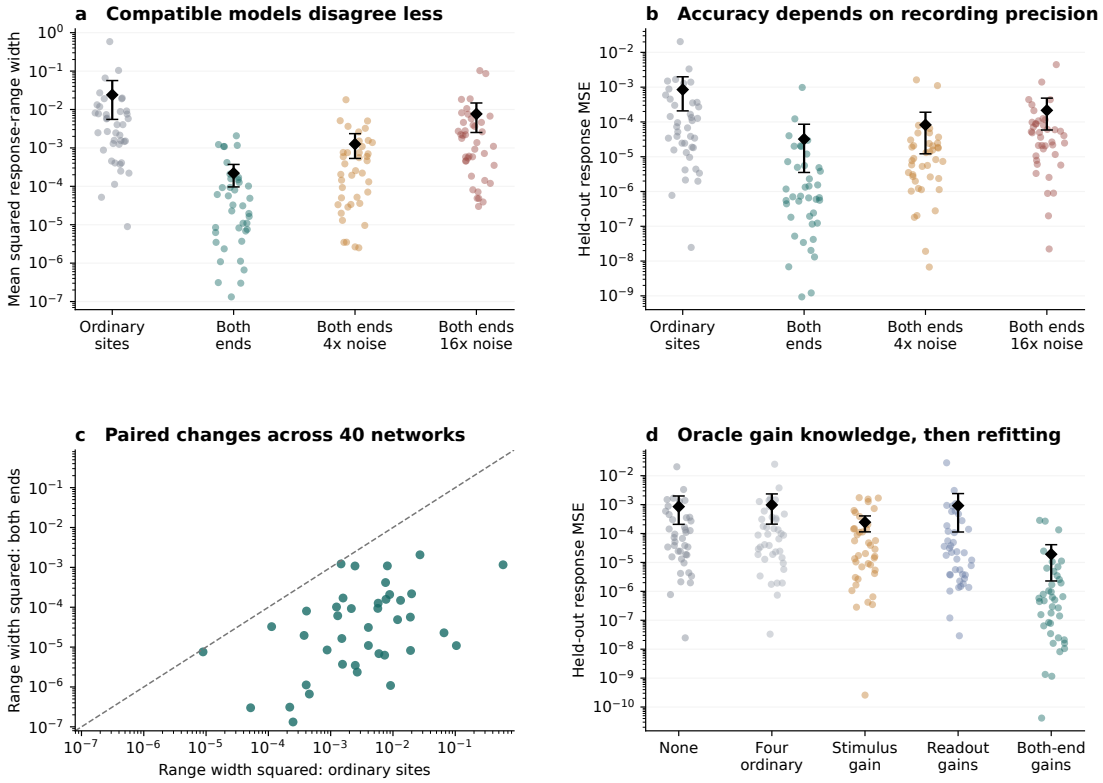
An information intervention distinguished endpoint-gain information from generic parameter knowledge. Keeping the five nonendpoint observations fixed, we revealed either the stimulation gain, all three readout gains, all four endpoint gains or four unrecorded nonendpoint gains, and refitted all remaining gains. Revealing both ends changed response MSE by -98.03% (95% CI -99.26% to -92.77%) relative to the four-ordinary-gain control, with improvement in 39/40 networks. The changes relative to knowing only the stimulation gain or only the readout gains were -92.19% (95% CI -98.98% to -79.73%) and -97.91% (95% CI -99.03% to -91.54%). These oracle controls are mechanism diagnostics; no revealed gain was available to a deployable selector.

Together, the feasible-response, precision and parameter-information comparisons support a specific mechanism in this model family: endpoint calibration constrains uncertainty relevant to the intended intervention response. This complements the earlier observation that accurate response prediction need not coincide with accurate recovery of each individual gain. It does not require interpreting descriptive error cancellation as an independent mechanism test.

2.8 Approximation diagnostics constrain the interpretation

Development-only diagnostics measured local Taylor remainders and selection overlap. Mean calibration remainders relative to parameter-induced activity changes were approximately 0.179 under matched uncertainty, 0.334 under broad uncertainty and 0.252 under hidden bias shift. Broader uncertainty changes signal scale and parameter variability as well as nonlinear behaviour, so these differences do not isolate a curvature effect. Local and ensemble sets frequently overlapped. When sets were identical, the fitted data and response predictions were identical by construction, limiting any decision benefit of a changed risk value in those cases.

The matched and hidden-shift conditions used the same public prior and therefore produced exactly the same prospective designs, despite different true responses. This is an information boundary: without a preliminary measurement, the selector cannot infer an undisclosed physiological shift. In the first frozen evaluation, approximate 95% local MAP response intervals for the nominal scalar estimator covered the truth with proportions 0.930, 0.919 and 0.793 for matched, broad and shifted conditions, respectively, at excluded endpoints and M=3. These diagnostics identify limitations of the current information and approximation contract; they do not establish a single biological mechanism for the observed errors.



Dots: independent networks. Diamonds: arithmetic means and 95% bootstrap intervals.
Response ranges use a common MAP-objective allowance of 2; they are numerical feasible spans, not confidence intervals.

Fig. 7 | Endpoint information and response ambiguity. a, Mean squared numerical response-range width under the same MAP-objective allowance. **b,** Held-out response MSE under the same allocations and endpoint-noise manipulations. **c,** Paired range changes across the 40 independent networks. **d,** Oracle gain knowledge followed by refitting the remaining parameters, using five nonendpoint recordings in every condition. Dots denote networks; diamonds and error bars denote arithmetic means and 95% network-bootstrap intervals. Both-versus-neither range width is the primary comparison; precision and oracle contrasts are secondary. Numerical feasible response ranges and bootstrap intervals over networks answer different questions.

3 Discussion

Joint calibration of the stimulation and readout sites is the central positive result. Its prediction benefit replicated across independent networks, and the further response-range and information-intervention experiments connect that benefit to constraints on the intended response. This supplies a common explanation for the allocation experiments and a useful baseline for evaluating risk-based selection. The original criterion comparisons and the additional indirect-recording study ask whether a more elaborate approximation adds value after the available information and fitting procedure are specified.

A non-significant interval is more informative when read as an approximate bound than as a verdict [26, 27]. In the original 12-unit primary study, the shared three-response 97.5% interval (−2.95% to +11.51%) does not extend below a mean-MSE reduction of about 3%, while it does extend to a deterioration of about 12%. For the single-response objective the interval (−18.48% to +12.12%) is too wide to exclude either a substantial gain or a substantial loss, and that comparison is where additional independent networks would be most informative. These readings are approximate: 13 and 26 of the 40 network differences are exactly zero and the rest are heavy-tailed, so percentile bootstrap intervals are not sharp, and no equivalence margin was specified in advance.

The study builds on recording-site optimisation in connectome-constrained models [9] and on goal-oriented experimental design [20, 22]. Bayesian trajectory-oriented design has also been compared with

Fisher-based design in dynamical models [28], and optimisation and identifiability diagnostics have an established literature [29–31]. The specific contribution here is the controlled distinction among risk approximation, endpoint allocation and shared fitting for a held-out neural intervention response. The independent endpoint result supplies a positive empirical finding, while the stronger baseline determines how much of the apparent design benefit requires sophisticated selection.

The independent mechanism follow-up links endpoint calibration to a narrower set of responses compatible with a common fitting criterion. The precision manipulation tests the information carried by those recordings, while the gain-information intervention identifies a task-relevant parameter block. Together they provide direct support beyond the earlier descriptive gain/dynamics error decomposition. Prediction profiles and prediction-guided experimental design already have established foundations [29, 30, 32]; the present contribution is the controlled evaluation of this mechanism for the specified neural intervention task. Numerical range searches and multiple starts remain approximations to a globally explored parameter space.

3.1 Limitations

The 12-unit and 100-unit families serve different purposes: the former supplies the original factorial and forecast studies, while the latter supplies independent confirmation and focused allocation experiments. The uncertainty-width and fitting analyses reuse the 40-network confirmation cohort. The 40-network endpoint validation uses disjoint networks. The response-range follow-up adds 40 further networks and the indirect-recording comparison adds a separate 40. Their sample sizes and primary comparisons were fixed locally before their respective results; precision and oracle contrasts are secondary. Numerical response ranges use a MAP-objective threshold rather than a calibrated frequentist coverage threshold.

The study assumes known numerical connectivity, a correct observation equation, known initial states and a specified stimulus. There is no process noise, plasticity or biological nonstationarity. The endpoint result concerns one rank-two model family and a three-readout task; its practical use depends on whether those sites can be recorded during calibration. Local risk, finite prior ensembles, greedy selection and MAP optimisation remain approximations. Multiple starts do not guarantee global optimisation. Joint neural recordings and perturbations would be needed to establish biological validity, and a further neuromuscular layer would be needed for mechanical predictions in biohybrid robots [2, 4].

The supported mechanism is conditional on the model, calibration inputs and specified response. Its extension to different network structures or physiological misspecification remains an external-validity question. The additional indirect-recording comparison evaluates a concrete computational alternative within the same task, with its performance reported separately from any claim of methodological novelty.

3.2 Implications for recording-design practice

An informative baseline should reflect the known stimulus and readouts. When endpoints are available during calibration, comparing a risk-based selector only with unrestricted random sampling can substantially overstate the value of sophisticated selection. The endpoint baseline offers a direct, inexpensive comparison. When endpoints are inaccessible, the criteria must earn their benefit through indirect recordings, which remains uncertain in the present primary confirmation.

4 Conclusions

The study connects recording choice, nonlinear response prediction and predictive-risk estimation under a common information budget. Joint calibration of the stimulation and readout sites provides a replicated positive allocation result. Independent response-range and information-intervention experiments support a mechanism in which endpoint information constrains the intended response, even when individual physiological parameters remain uncertain. Risk-based and covariance-based selectors are evaluated against this information context; their incremental performance is distinct from the value of recording the relevant sites. The resulting contribution is a controlled empirical account of calibration information and intervention prediction in the tested network family.

5 Methods

5.1 Notation

Symbol	Meaning
J, N	known connectivity matrix and number of units
$\theta = (\eta, b)$	uncertain log-gains and biases
M	recording budget (number of recorded neurons)
q, Q	specified target response and the three-response vector
F, H	whitened observation sensitivities and local information matrix
R	local delta-method risk of a recording set
v, e	forecast of the mean squared error and the realised prediction error
S(v, e)	forecast score $\log v + e^2/v$
LS, ES, LV, EV	nominal and ensemble criteria for the single-response and three-response objectives

5.2 Model, measurements and response

The primary model is

$$\dot{x} = -x + Jr + u, \qquad r_i = \exp(\eta_i)\,\mathrm{softplus}(x_i + b_i). \tag{1}$$

N=12, x(0)=0 and J is known. The independent synthetic J matrices have 30% sparse connectivity, zero diagonal, column-wise excitatory/inhibitory signs, and spectral norm 0.65. The uncertain vector is θ=(η,b). The observation is the rate r, not the internal state x. Time and activity are dimensionless.

Two known calibration inputs apply sinusoidal probes to three channels. Euler integration uses a step of 0.05 over six time units. Sampling every fourth step yields 60 observations per recorded neuron across the two calibration trials. Measurement noise is independent Gaussian with standard deviation 0.05.

A separate fixed stimulus of amplitude 0.6 acts during time 1–2. Three readouts are selected from structurally reachable nodes using J before physiological parameters or measurement noise are generated. The target family is

$$q_k(\theta) = \int_0^6 [r_{j_k}(t;\theta,u_*) - r_{j_k}(t;\theta,0)]dt. \tag{2}$$

The primary implementation evaluates this integral by the matching step-weighted (left-point) sum with dt=0.05. Q consists of these three responses. The vector criterion and primary raw-error evaluation use equal weights over the same target family. Scalar designs specialise for one known target; vector designs use one configuration for all three. No observations from separately selected scalar configurations are pooled.

Recording budgets are M=3 and M=6. The candidate pool either includes all neurons or excludes the stimulation site and all three readouts. The latter pool is shared by every target and method.

5.3 A factorial comparison of existing criteria

In whitened parameter coordinates, the local information matrix for a recording set M at parameter particle ℓ is

$$H_{M,\ell} = I + \sum_{i\in M} F_{i,\ell}^T F_{i,\ell}, \tag{3}$$

where F includes the observation derivatives, prior parameter scale and noise standardisation. With a_k,ℓ denoting the whitened target gradient, the scalar criterion is

$$R_{k,\ell}(M) = a_{k,\ell}^T H_{M,\ell}^{-1} a_{k,\ell}. \tag{4}$$

The vector criterion averages these risks over the three readouts. Local designs evaluate at μ=0; ensemble designs average the individually solved risks over 16 common prior draws. Averaging information before solving would define a different criterion and is not used.

Design	Predictive objective	Parameter representation
LS	One specified response	Nominal parameter
ES	One specified response	Common prior ensemble
LV	Three-response vector	Nominal parameter
EV	Three-response vector	Common prior ensemble
Activity	Full calibration activity	Nominal parameter
Random	No risk objective	Five feasible random sets

All risk-based methods use the same greedy forward selection and deterministic tie convention. These are greedy-selected sets, not certified global optima. Random permutations give nested budgets. The design API receives only the public model, prior, inputs and task. It receives neither generating parameters nor measured evaluation outcomes. There is no uncounted all-neuron pretraining or preliminary dataset.

The ensemble criterion is a **prior average of local delta-method risk**. It does not integrate future measurements and nonlinear posteriors, establish a variance upper bound, or capture all posterior modes.

5.4 Common inference and paired evaluation

After selection, all methods minimise the same Gaussian MAP objective, using analytic sensitivities and the prior mean as the initial parameter. The solver runs at most 80 function evaluations with identical tolerances. A finite final iterate is retained and flagged if convergence is not reached. Exceptions or nonfinite predictions trigger a flagged prior-mean fallback. No instance is silently dropped.

Within each network, true parameter vector and noise repetition, methods receive their selected subsets of one common measurement-noise field. The full field is internal to the generator. Identical recorded datasets receive the same MAP objective, prior-mean starting point, common noise field, solver settings and deterministic solver, so identical sets give identical fits and predictions by construction; they may also reuse a cached fitted solution; this changes compute accounting, not information access. Responses are predicted by nonlinear simulation at the fitted parameter. A passive input not used for fitting gives an additional activity-prediction diagnostic.

5.5 Development and the first frozen evaluation

Stage	Independent J	Physiological truths/J	Noise repeats/truth	Readouts/J	Conditions
Development	8	2	2	3	Matched, broad, bias shift, gain-only, bias-only
First frozen evaluation	40	3	3	3	Matched, broad, bias shift

Matched and broad conditions use correctly specified Gaussian priors and generating standard deviations 0.18 and 0.35 respectively. Bias shift retains the design prior at 0.18 while adding 0.5 to the hidden true bias mean. The method is not informed of this shift. Development gain-only and bias-only conditions fix the other parameter group. Increasing a prior scale changes parameter variability, signal scale and nonlinearity together; it is not an isolated curvature intervention.

Connectivity, truth, noise, design ensembles, task-scale Monte Carlo, random sets and analysis use separate seed namespaces. The first evaluation uses network IDs 5000–5039, disjoint from development. The execution protocol and computational source were locally frozen before outcomes were inspected. This is not external preregistration.

5.6 Estimands and uncertainty

For each J, squared response errors are averaged across all prescribed truths, noises, targets and random-design repetitions. Network risks are then averaged equally across J. RMSE is the square root of the overall mean squared error. Repeated targets or noises do not create additional independent networks.

Interval key: 97.5% intervals belong to the two prespecified first-study comparisons, the 95% interval to the single prespecified forecast comparison, and all other intervals are pointwise 95% and exploratory.

Two prespecified primary comparisons, ES−LS and EV−LV, concern broad uncertainty, excluded endpoints and M=3. The paired network-cluster percentile bootstrap [33, 34] uses 10,000 resamples of whole networks, so repeated targets and noises do not count as independent. Each primary comparison receives a 97.5% interval as a two-comparison Bonferroni interval arrangement. Other strata, task-specialisation comparisons and factorial interactions have exploratory pointwise 95% intervals. We do not infer equivalence from an interval crossing zero [26, 27].

The primary scale is raw response MSE. A secondary scale uses an independent common 64-draw nonlinear prior-Monte-Carlo variance. It is not method-specific. Prescribed targets below a frozen variance floor of 10^{-16} retain raw results; a complete standardised estimand is NA rather than silently excluding them. Local MAP intervals are approximate diagnostics, not guaranteed confidence intervals.

5.7 Separately specified risk-forecast follow-up

The first study distinguishes selection gains from apparent risk-estimation differences only exploratorily. A subsequent protocol therefore uses new network IDs 6000–6039 and fixes the recording policies to LS and LV. The point predictions and selected data are identical for the two risk forecasts being compared. The forecasts are the nominal and ensemble local risks evaluated at that same recording set.

The risk-forecast score is

$$S(v, e) = \log v + e^2/v, \qquad v = \max(\hat{R}, 10^{-24}). \tag{5}$$

For positive v, its conditional expectation is minimised at v=E[e^2]. This makes it a consistent score for a positive MSE forecast; it does not establish a full predictive distribution or require the errors to be Gaussian. The log term penalises indiscriminate inflation. This is a use of established scoring principles [35, 36], not a new score. The forecast-evaluation literature distinguishes calibration from sharpness [37] and treats calibration as separate from point-prediction accuracy [38]; the nominal-versus-ensemble comparison here mixes both, applies to a mean-squared-error forecast rather than a full predictive distribution, and does not test calibration separately.

The follow-up has one frozen primary comparison: ensemble minus nominal risk-forecast score for the LS recording policy under broad uncertainty, excluded endpoints and M=3. It uses a paired network-cluster 95% interval. Matched/shift conditions and LV are secondary. These results do not redefine the first study's primary hypotheses.

5.8 Numerical and published-model checks

The implementation is checked against an independent NumPy model, analytic finite differences, nonordered observation subsets, a linear bias-cancellation null, and exact small-subset enumeration. A separate nondegenerate linear-Gaussian positive control uses constant observation/target Jacobians, independent prior truths and noise, and exact linear MAP predictions. It checks both nominal–ensemble equivalence and the link between $a^T Ca$ and actual MSE. The original SVD criterion is independently calibrated in its noiseless linear internal-state setting against Monte Carlo reconstruction, rather than mislabelled as a noisy-rate baseline reproduction.

The N=100 model-family extension uses the published rank-two tanh distribution, raw signed gains g~N(1,0.9^2), zero bias and τ=0.5 [9]. Its inputs, known initial state, measurement noise and inference are explicit adaptations. Negative gains are retained as mathematical parameters. The coarse Euler implementation did not pass continuous-time convergence checks. The refined implementation uses x=v+Lz, analytic filtering of the known held input, and RK4 on the reduced state and sensitivities. Responses and their complete gradients are integrated at RK4 stages. It is compared with dense RK4 and finer substeps before execution. The initial pilot used eight new networks, two truths, two noise repetitions, three targets, budgets 5/10 and both pools. The prior ensemble has 32 draws and MAP permits 200 evaluations. These settings are frozen before formal extension outcomes, following numerical and timing checks. All initial pilot comparisons are exploratory. This extension changes both model family and N, so it is not a controlled estimate of an N-only scaling law.

5.9 Independent confirmation and focused follow-ups

The 100-unit confirmation used 40 new networks (IDs 20000–20039), two independent physiological draws, two noise draws and three fixed readouts per network, both recording pools and budgets five and ten. The primary contrast was ensemble versus nominal selection for the shared objective, with endpoints excluded and budget five. The original model, integration, noise, 32-particle design ensemble and common MAP rule were retained. The sample size and primary contrast were locally frozen before inspecting outcomes. All other confirmation contrasts use exploratory 95% network-paired bootstrap intervals.

The matched-width analysis reused these networks with prior and truth standard deviations 0.3, 0.6 and 0.9, scaling common normal draws. First-step rank correlations used a common empty selection prefix. Component diagnostics held either target sensitivities or observation information at the nominal values. The fitting follow-up kept the selected sets fixed and fitted from zero, two draws scaled to 0.25 prior standard deviations and two full prior draws. Starts were common to both selectors and data repetitions within a network (seed stream 91631). The fit with the smallest unchanged MAP objective supplied the response prediction. True parameters were used only for diagnostics: a returned objective above the objective at the generating parameters certifies a suboptimal returned fit, but the reverse does not certify optimality.

The endpoint experiment fixed five recordings and used four allocations: five nonendpoints; the stimulus site plus four nonendpoints; the three readouts plus two nonendpoints; and all four endpoints plus one nonendpoint. For each of three allocation repeats, a common random permutation of nonendpoints supplied the remaining sites (streams 1601–1603). Thus policies had identical total measurement counts and shared truth and observation-noise draws. The original zero-start MAP estimator was used throughout. The first cohort reused networks 20000–20039. A separate prospective local protocol fixed both-versus-neither as the primary comparison on 40 new networks (IDs 30000–30039); both-versus-stimulus-only and both-versus-readouts-only were secondary comparisons. All 40 planned networks were retained in each cohort.

Response squared errors were averaged over targets, truth draws, noise draws and allocation repeats within each network. Contrasts are ratios of means across networks minus one, with 10,000 paired network bootstrap resamples. Evaluation rows and fit restarts are not treated as independent observations. Local freezes and code hashes accompany the outputs; they are not public preregistrations.

For the local information diagnostic, parameters were prior-whitened and the nominal Gaussian covariance was C=(I+F_S^{T}F_S)$^{-1}$. Marginal variances were compared with the unit prior variance. If an endpoint parameter block E were known exactly, the conditional covariance would be C−C_:E C_EE^{-1} C_E:. The reduction in trace(GCG^{T})/3 quantifies the local risk removable by that additional knowledge, including parameter correlations. This is an approximation diagnostic rather than an empirical variance estimate. A separate nonlinear oracle check substituted true endpoint gains into fitted parameter vectors, held all other fitted values fixed and recomputed the response; oracle outcomes never entered policy or fit selection.

Task specificity was evaluated by simulating the already fitted models under the original other-input test traces and comparing whole-network activity with the teacher. For the descriptive response-error decomposition, write $q_i=g_iL_i$, where L_i is the time-integrated difference of tanh states under stimulated and control trajectories. Then $\hat{q}_i - q_i = (\hat{g}_i - g_i)L_i + \hat{g}_i(\hat{L}_i - L_i)$. We report the mean square of each term and twice their mean product. The cancellation fraction is minus this cross term divided by the sum of the two squared terms, aggregated over the same records as the response MSE. The sum agreed with the response MSE for every saved fit; the decomposition describes compensation in fitted models rather than a separate intervention on biological dynamics.

5.10 Response ambiguity and gain-information interventions

The mechanism study used network IDs 40000–40039, one prespecified physiological draw and noise draw per network, and the same three readouts. Five common starts comprised the nominal mean, two draws at one quarter of the prior standard deviation and two full prior draws (stream 93631). Fitting retained the smallest MAP objective with a 300-evaluation cap. The response-range problem maximised or minimised each target over the full nonlinear MAP sublevel set with allowance delta=2. SLSQP used two starts per direction, a 160-iteration cap, analytic sensitivities and invertible local Gauss-Newton coordinates. The quadratic prior makes the original finite coordinate bounds redundant. Every retained extreme was checked against the original objective; the final numerical envelope also

included every known feasible multistart fit, preserving the original search output separately; discovering a lower objective triggered refitting and a repeated range search. Value-only line searches and the coordinate transformation preserved the scientific objective, feasible region, samples and comparisons; their numerical checks and timing amendments were logged before completed validation outputs.

The primary range metric averaged squared marginal response widths across the three targets, then across networks. Secondary experiments multiplied endpoint-noise standard deviations by four and sixteen while keeping their standardised random draws paired. The first eight IDs were fixed for a delta=5 threshold check. The joint-endpoint versus ordinary-site range-metric change was -99.77% (95% CI -99.85% to -93.61%) in that check. Known feasible delta=2 extrema were retained in the wider search envelope. All retained extrema were reevaluated with twice the integration resolution; the maximum response change was 0.00231 of the corresponding range width. The primary comparison retained 479 successfully optimized boundary points and one MAP point. Two of 480 retained points in the noise conditions and two of 96 points in the wider-threshold check had unsuccessful optimizer flags but satisfied the numerical feasibility tolerance. The finer integration changed the primary range-metric effect from -99.08477% to -99.08468%. Ranges are feasible numerical spans, not guarantees of global extrema or calibrated confidence intervals.

Oracle information controls used the same five nonendpoint recordings and fixed the revealed gains to their generating values before refitting every remaining gain. The matched ordinary-gain control revealed four gains outside both the endpoints and the recorded set (stream 93931). These diagnostic gains never entered a recording selector. Statistical comparisons resampled complete networks 10,000 times, keeping all within-network contrasts paired.

5.11 Indirect-recording comparison using joint prior simulations

A separate locally frozen comparison used networks 45000–45039, two physiological draws, two noise draws and three targets per network. All selectors chose five nonendpoint sites and used the same five-start nonlinear MAP estimator. In addition to nominal and prior-averaged local-risk selection, a joint covariance criterion estimated observation–response dependence from 32 prior simulations. A held-out-prior variant evaluated the resulting linear predictor on 32 independent prior simulations and added the analytic contribution of observation noise to its squared error. Greedy choices minimised this surrogate risk. No realised calibration or held-out stimulus response was available during selection. This construction draws on established goal-oriented and simulation-based design ideas [20, 25, 39].

The held-out-prior versus nominal MSE contrast was primary. Covariance, ensemble and random comparisons were secondary. A retrospective optimization audit identified 11 truth/noise datasets in which at least one selected fit had an objective above that at the generating parameters. All five policies for each flagged dataset received the same eight additional prior-generated starts (stream 94031), plus refinement of the old best fit. The generating parameters flagged the diagnostic but were never starting points or selection criteria. These sensitivity results are reported separately from the frozen comparisons. The final response estimates always came from the common nonlinear fit, so the comparison evaluates whether the surrogate selects useful recordings for that estimator. Numerical tests checked the dual covariance formula against direct observation-space calculations and verified the response and objective derivatives.

5.12 Study accounting and availability

Core runtime is recorded separately from environment preparation, compilation, literature review and figure production. Dataset reuse avoids repeated optimisation of identical selected observations; evaluation rows are not independent sample counts.

The independent 100-unit confirmation added 40 networks, 15,360 scalar response records and 7,112 actual MAP calls. The matched-width diagnostics reused those networks (11,520 records; 3,296 calls). The focused endpoint experiment reused the same 40 networks (1,920 joint-response fit records; 1,916 calls), and independent endpoint validation added 40 new networks (1,920 records; 1,908 calls). The common five-start fitting analysis reused 320 datasets and made 1,600 fit calls. Two endpoint-validation calls reached the evaluation cap with finite outputs, which were retained. Thirty-seven alternative starts in the fitting follow-up reached the cap; all selected best fits terminated successfully. Before the final follow-ups, the 100-unit evidence comprised eight pilot networks, forty criterion-confirmation networks

Table 3 | Study accounting

Stage	Independent J	Truth × noise × readouts per J	Evaluation records	Actual MAP calls	Core runtime (s)
Development	8	2 × 2 × 3	19,200	5,696	25.4
Frozen recording evaluation	40	3 × 3 × 3	129,600	38,610	356.0
Independent forecast study	40 new	3 × 3 × 3	6,480	3,006	17.5
Exploratory 100-unit family	8 new	2 × 2 × 3	3,072	1,436	693.1

Connectivity matrices are the independent sampling units. Truth, measurement noise and readout repeats are nested within J. Timings are core experiment wall times on the study Mac and are not cross-platform performance guarantees.

and forty additional endpoint-validation networks. The final cohorts are accounted for below; none is used to replace an earlier frozen comparison.

The companion package contains the independent implementation, dependency lock, actual execution/source freezes, per-instance outputs, numerical controls, network-level analysis and figure source data. MIASE, DOME and MDAR inform reporting [40–42]. No new human or animal measurements were collected. Source equations were independently implemented; unlicensed upstream code is not redistributed. All freezes were local records rather than external preregistrations.

The response-range study and the indirect-recording study each add forty independent 100-unit networks, bringing the total to 168: eight pilot networks and four distinct forty-network cohorts. Oracle and threshold diagnostics reuse the response-range cohort. These cohorts are not pooled across differing estimators or experimental questions.

6 Data and code availability

The original frozen analysis tables and figure source data remain in the collaborator's research package. The repository workspace `work/recording_extensions` contains the independent confirmation and matched-width experiments; its `round2` directory contains the endpoint cohorts, common fitting follow-up, source code, protocols, parameter-level outputs and analysis tables. The round3 directory adds the two new independent cohorts, locally frozen protocols, full numerical response searches, gain-information interventions, indirect-recording criteria, numerical audits and source figures. The same workspace contains all three extensions, their dependency lock, and integrity checks. Upstream code and weights from the published model family are not redistributed.

7 Author contributions

J.D. and S.C. contributed equally to this work, including conceptualisation, methodology, investigation, visualisation and writing.

8 Competing interests and funding

The authors declare no competing interests. This work received no specific funding. No human or animal measurements were collected.

References

[1] Kim, Y. *et al.* Remote control of muscle-driven miniature robots with battery-free wireless optoelectronics. *Science Robotics* **8**, eadd1053 (2023).

[2] Min, H. *et al.* Optogenetic neuromuscular actuation of a miniature electronic biohybrid robot. *Science Robotics* **10**, eadu5830 (2025).

[3] Ricotti, L. *et al.* Biohybrid actuators for robotics: A review of devices actuated by living cells. *Science Robotics* **2**, eaaq0495 (2017).

[4] Aydin, O. *et al.* Neuromuscular actuation of biohybrid motile bots. *Proceedings of the National Academy of Sciences* **116**, 19841–19847 (2019).

[5] Wang, X. *et al.* Biomimetic Visual Information Spatiotemporal Encoding Method for In Vitro Biological Neural Networks. *Biomimetics* **10**, 359 (2025).

[6] DeMarse, T. B., Wagenaar, D. A., Blau, A. W. & Potter, S. M. The Neurally Controlled Animat: Biological Brains Acting with Simulated Bodies. *Autonomous Robots* **11**, 305–310 (2001).

[7] Prinz, A. A., Bucher, D. & Marder, E. Similar network activity from disparate circuit parameters. *Nature Neuroscience* **7**, 1345–1352 (2004).

[8] Gutenkunst, R. N. *et al.* Universally Sloppy Parameter Sensitivities in Systems Biology Models. *PLoS Computational Biology* **3**, e189 (2007).

[9] Beiran, M. & Litwin-Kumar, A. Prediction of neural activity in connectome-constrained recurrent networks. *Nature Neuroscience* **28**, 2561–2574 (2025).

[10] Lappalainen, J. K. *et al.* Connectome-constrained networks predict neural activity across the fly visual system. *Nature* **634**, 1132–1140 (2024).

[11] Lal, T. N. *et al.* Support Vector Channel Selection in BCI. *IEEE Transactions on Biomedical Engineering* **51**, 1003–1010 (2004).

[12] Jun, J. J. *et al.* Fully integrated silicon probes for high-density recording of neural activity. *Nature* **551**, 232–236 (2017).

[13] Haber, A., Molnar, F. & Motter, A. E. State Observation and Sensor Selection for Nonlinear Networks. *IEEE Transactions on Control of Network Systems* **5**, 694–708 (2018).

[14] Krause, A., Singh, A. & Guestrin, C. Near-Optimal Sensor Placements in Gaussian Processes: Theory, Efficient Algorithms and Empirical Studies. *Journal of Machine Learning Research* **9**, 235–284 (2008). URL https://jmlr.org/papers/v9/krause08a.html.

[15] Sourmpis, C., Petersen, C. C., Gerstner, W. & Bellec, G. Biologically informed cortical models predict optogenetic perturbations. *eLife* **14**, RP106827 (2026).

[16] Ogino, M., Sekizawa, D., Kitazono, J. & Oizumi, M. Designing optimal perturbation inputs for system identification in neuroscience. eLife Reviewed Preprint, version 1 (2026).

[17] Lewi, J., Butera, R. & Paninski, L. Sequential Optimal Design of Neurophysiology Experiments. *Neural Computation* **21**, 619–687 (2009).

[18] Chaloner, K. & Verdinelli, I. Bayesian Experimental Design: A Review. *Statistical Science* **10**, 273–304 (1995).

[19] Lindley, D. V. On a Measure of the Information Provided by an Experiment. *The Annals of Mathematical Statistics* **27**, 986–1005 (1956).

[20] Attia, A., Alexanderian, A. & Saibaba, A. K. Goal-oriented optimal design of experiments for large-scale Bayesian linear inverse problems. *Inverse Problems* **34**, 095009 (2018).

[21] Alexanderian, A., Petra, N., Stadler, G. & Ghattas, O. A Fast and Scalable Method for A-Optimal Design of Experiments for Infinite-dimensional Bayesian Nonlinear Inverse Problems. *SIAM Journal on Scientific Computing* **38**, A243–A272 (2016).

[22] Zhong, S., Shen, W., Catanach, T. & Huan, X. Goal-Oriented Bayesian Optimal Experimental Design for Nonlinear Models Using Markov Chain Monte Carlo. *SIAM/ASA Journal on Uncertainty Quantification* **14**, 19–47 (2026).

[23] Ryan, E. G., Drovandi, C. C., McGree, J. M. & Pettitt, A. N. A Review of Modern Computational Algorithms for Bayesian Optimal Design. *International Statistical Review* **84**, 128–154 (2015).

[24] Rainforth, T., Foster, A., Ivanova, D. R. & Bickford Smith, F. Modern Bayesian Experimental Design. *Statistical Science* **39**, 100–114 (2024).

[25] Huan, X. & Marzouk, Y. M. Simulation-based optimal Bayesian experimental design for nonlinear systems. *Journal of Computational Physics* **232**, 288–317 (2013).

[26] Amrhein, V., Greenland, S. & McShane, B. Scientists rise up against statistical significance. *Nature* **567**, 305–307 (2019).

[27] Altman, D. G. & Bland, J. M. Statistics notes: Absence of evidence is not evidence of absence. *BMJ* **311**, 485 (1995).

[28] Weber, P., Kramer, A., Dingler, C. & Radde, N. Trajectory-oriented bayesian experiment design versus fisher a-optimal design: an in depth comparison study. *Bioinformatics* **28**, i535–i541 (2012). URL https://doi.org/10.1093/bioinformatics/bts377.

[29] Kreutz, C., Raue, A. & Timmer, J. Likelihood based observability analysis and confidence intervals for predictions of dynamic models. *BMC Systems Biology* **6**, 120 (2012).

[30] Raue, A. *et al.* Structural and practical identifiability analysis of partially observed dynamical models by exploiting the profile likelihood. *Bioinformatics* **25**, 1923–1929 (2009).

[31] Raue, A. *et al.* Lessons learned from quantitative dynamical modeling in systems biology. *PLOS ONE* **8**, e74335 (2013). URL https://doi.org/10.1371/journal.pone.0074335.

[32] Flassig, R. J., Migal, I., van der Zalm, E., Rihko-Struckmann, L. & Sundmacher, K. Rational selection of experimental readout and intervention sites for reducing uncertainties in computational model predictions. *BMC Bioinformatics* **16**, 13 (2015). URL https://doi.org/10.1186/s12859-014-0436-5.

[33] Efron, B. Bootstrap Methods: Another Look at the Jackknife. *The Annals of Statistics* **7**, 1–26 (1979).

[34] Field, C. A. & Welsh, A. H. Bootstrapping Clustered Data. *Journal of the Royal Statistical Society Series B: Statistical Methodology* **69**, 369–390 (2007).

[35] Gneiting, T. & Raftery, A. E. Strictly Proper Scoring Rules, Prediction, and Estimation. *Journal of the American Statistical Association* **102**, 359–378 (2007).

[36] Gneiting, T. & Katzfuss, M. Probabilistic Forecasting. *Annual Review of Statistics and Its Application* **1**, 125–151 (2014).

[37] Gneiting, T., Balabdaoui, F. & Raftery, A. E. Probabilistic Forecasts, Calibration and Sharpness. *Journal of the Royal Statistical Society Series B: Statistical Methodology* **69**, 243–268 (2007).

[38] Dawid, A. P. The Well-Calibrated Bayesian. *Journal of the American Statistical Association* **77**, 605–610 (1982).

[39] Thibaut, R. *et al.* Comparing well and geophysical data for temperature monitoring within a bayesian experimental design framework. *Water Resources Research* **58**, e2022WR033045 (2022). URL https://doi.org/10.1029/2022WR033045.

[40] Waltemath, D. *et al.* Minimum Information About a Simulation Experiment (MIASE). *PLoS Computational Biology* **7**, e1001122 (2011).

[41] Walsh, I. *et al.* DOME: recommendations for supervised machine learning validation in biology. *Nature Methods* **18**, 1122–1127 (2021).

[42] Macleod, M. *et al.* The MDAR (Materials Design Analysis Reporting) Framework for transparent reporting in the life sciences. *Proceedings of the National Academy of Sciences* **118**, e2103238118 (2021).

Supplementary figures

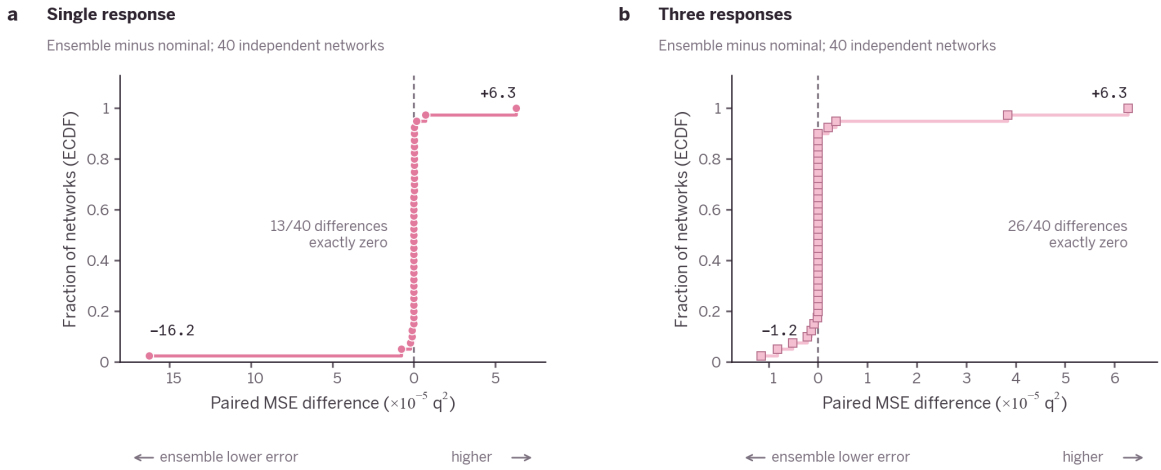


Fig. S1 | Complete paired primary differences. Each panel shows all 40 independent network-level ensemble-minus-nominal raw-MSE differences for the single-response (a) and three-response (b) comparison, including exact ties and the influential networks; the step curve is the empirical distribution function and every sorted value is drawn as a marker. The horizontal axis is in units of 10^{-5} q^2; negative values favour the ensemble. Extreme values (-16.2 and $+6.3$ in a; -1.2 and $+6.3$ in b) are labelled, not truncated. Exactly 13 (a) and 26 (b) of the 40 differences are zero. The distributions are descriptive; no interval is drawn, and the primary intervals remain those in Table 1 and Fig. 3. No equivalence margin was specified, and nothing here shows that the two selection rules perform alike.

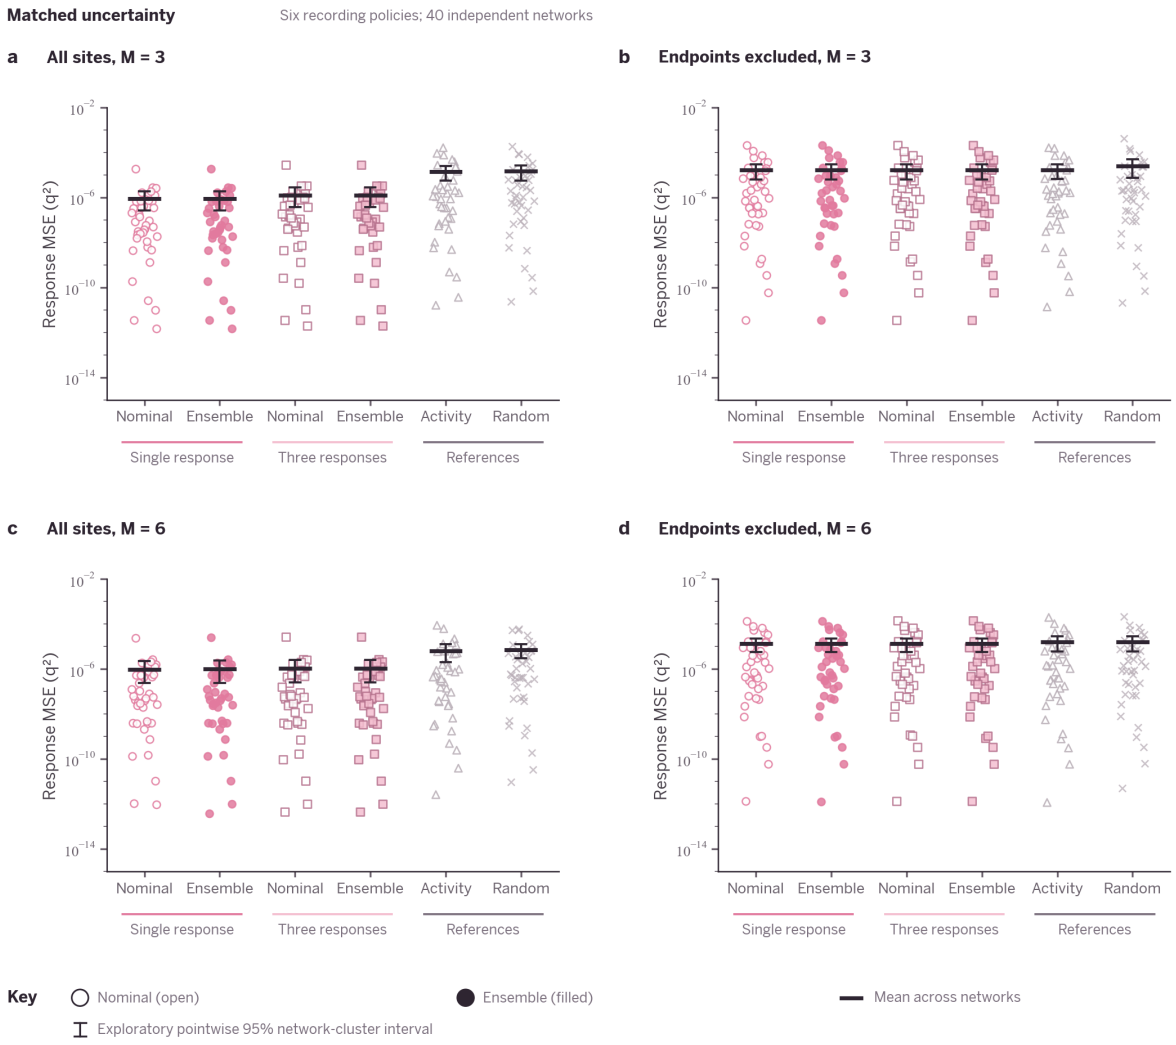


Fig. S2 | All recording conditions under matched uncertainty. Panels a–d show all-site/M = 3, endpoint-excluded/M = 3, all-site/M = 6 and endpoint-excluded/M = 6. For each of six policies (nominal and ensemble single-response, nominal and ensemble three-response, activity, random) every one of the 40 independent networks is shown as one point, with truth, noise and readout repeats averaged within network. Thick black bars are the aggregate means and thin black lines the exploratory pointwise 95% network-cluster intervals from the frozen summary table. Open markers denote nominal and filled markers ensemble selection; horizontal jitter only separates overlapping points. All panels share the logarithmic response-MSE axis (10^{-15} to 10^{-2}). No point or tail is removed. Intervals are not multiplicity-adjusted and no pairwise or equivalence claim is made.

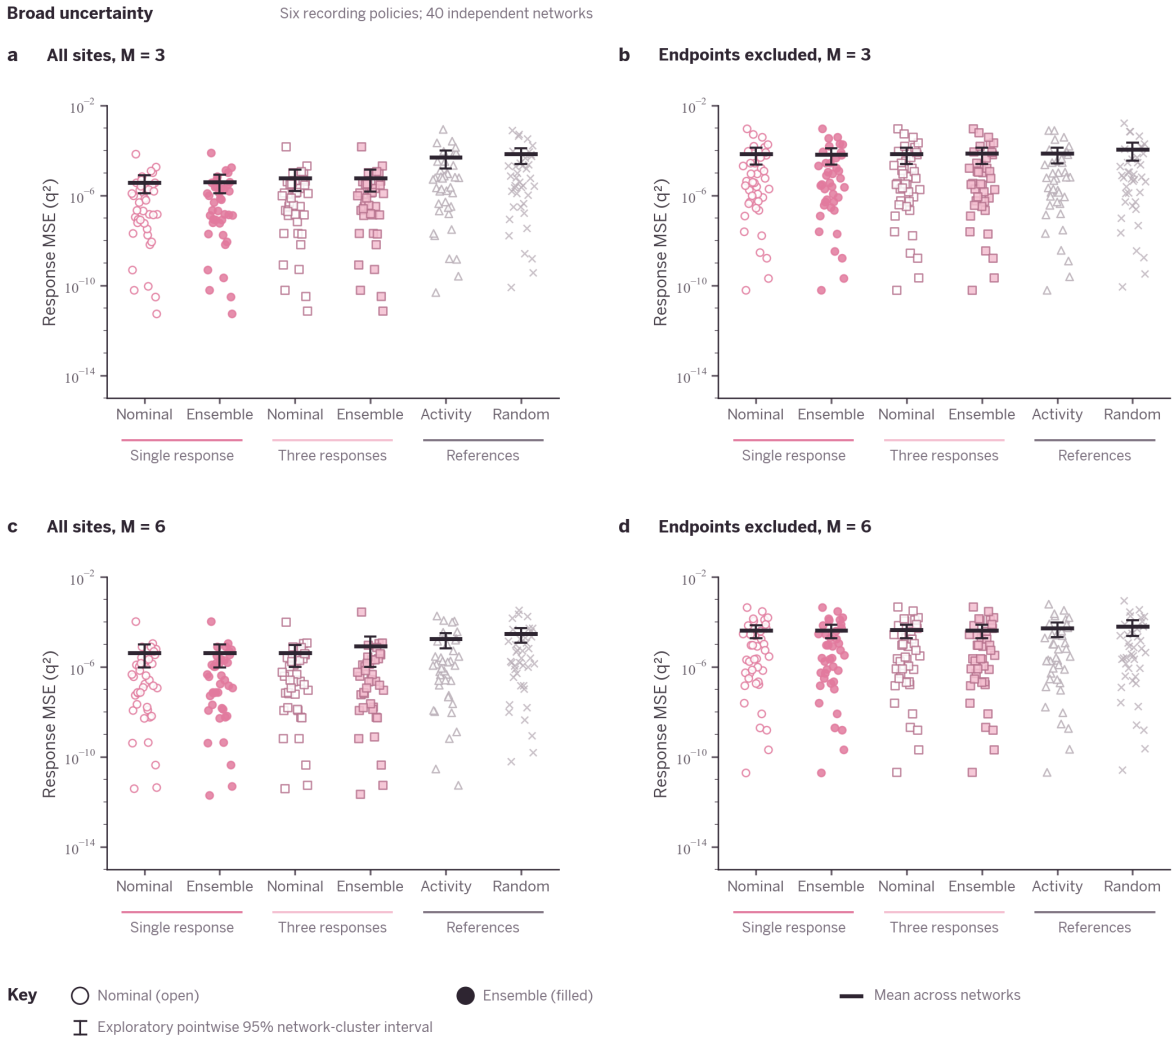


Fig. S3 | All recording conditions under broad uncertainty. As Supplementary Fig. 2 for the broad prior: panels a–d show all-site/M = 3, endpoint-excluded/M = 3, all-site/M = 6 and endpoint-excluded/M = 6; 40 independent networks per policy and panel; black bars and lines are aggregate means and exploratory pointwise 95% network-cluster intervals; shared logarithmic axes; open nominal, filled ensemble; jitter is horizontal only. The primary comparison in Fig. 3 corresponds to panel b. No method-specific tail is removed, and the panels do not support any claim of equivalence or superiority.

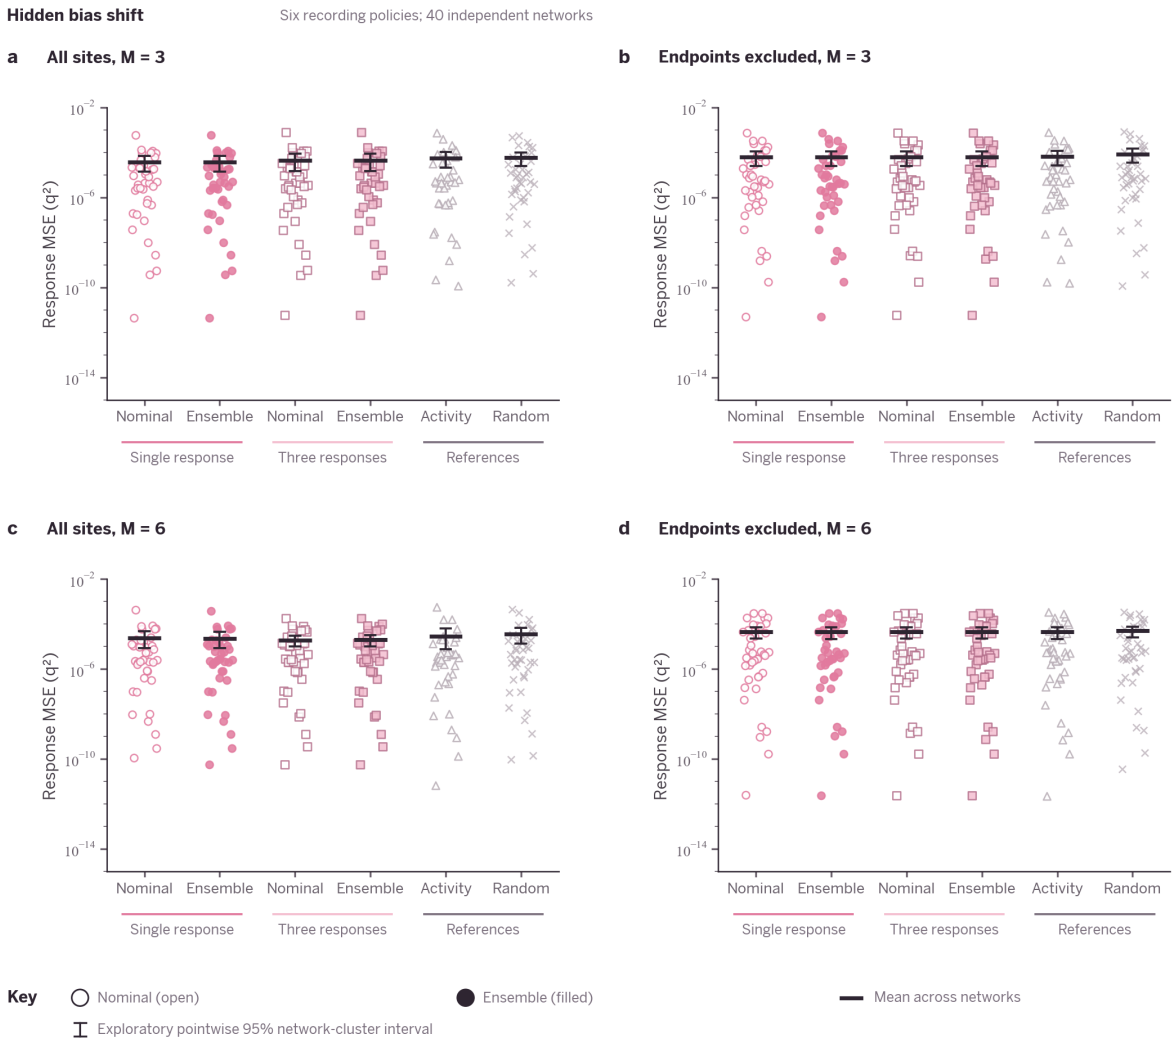


Fig. S4 | All recording conditions under hidden bias shift. As Supplementary Fig. 2 for the hidden-bias-shift condition: panels a–d show all-site/M = 3, endpoint-excluded/M = 3, all-site/M = 6 and endpoint-excluded/M = 6; 40 independent networks per policy and panel; black bars and lines are aggregate means and exploratory pointwise 95% network-cluster intervals; shared logarithmic axes; open nominal, filled ensemble; jitter is horizontal only. No data are filtered, and the intervals are descriptive and not multiplicity-adjusted.

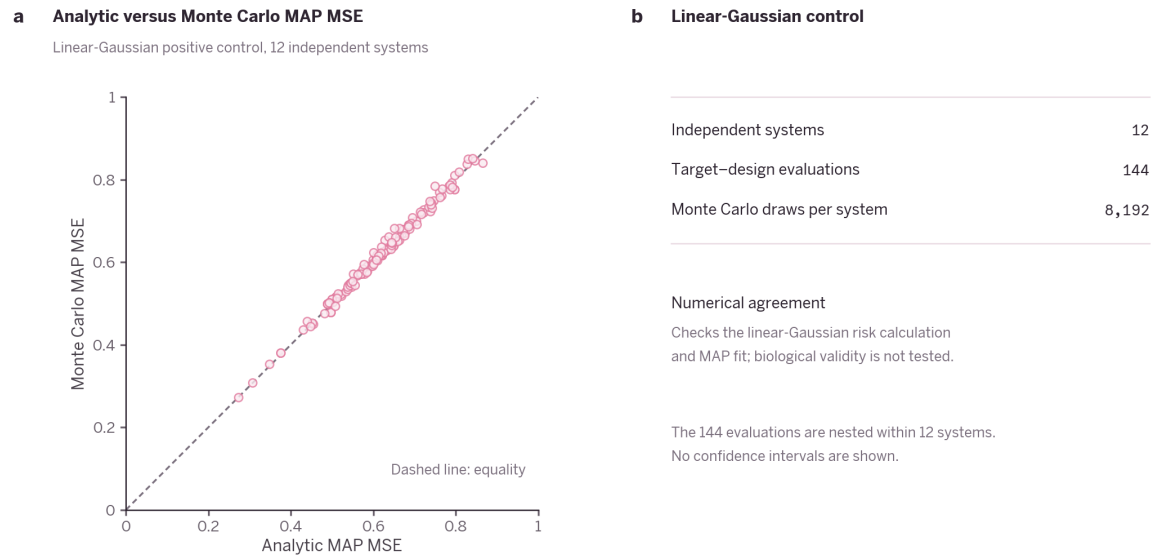


Fig. S5 | Nondegenerate linear-Gaussian positive control. All 144 target/design evaluations from 12 independent linear systems compare the analytic MAP MSE (horizontal) with the Monte Carlo MSE from 8,192 draws per system (vertical); the dashed identity line is the analytic reference. The independent unit is the system (n = 12); the 144 points are evaluations within systems, not independent neural networks, and no confidence intervals are displayed. Agreement validates this computational special case only; it does not validate biological assumptions.

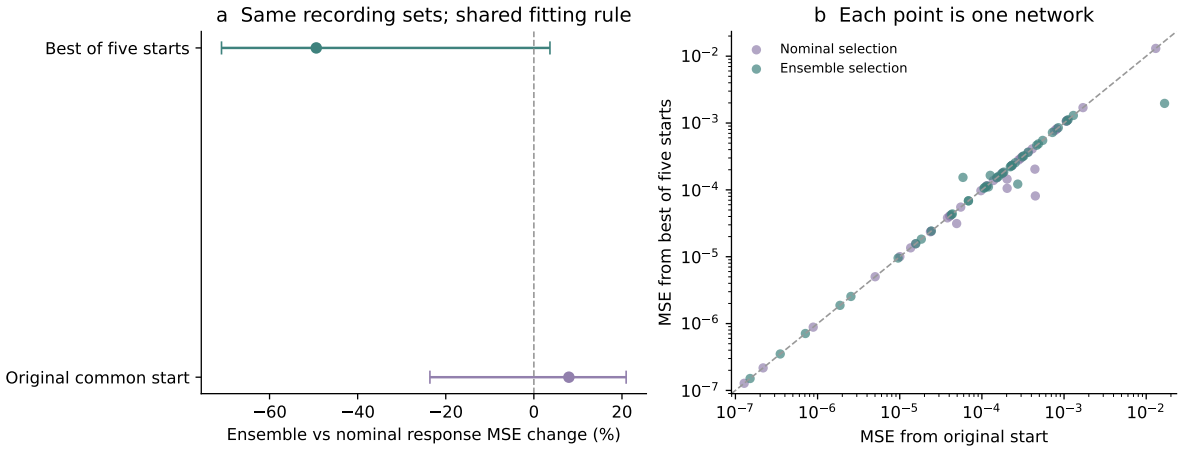


Fig. S6 | Fitting sensitivity in the primary 100-unit condition. Recording sets are fixed; both selectors receive the same five-start rule and retain the smallest MAP objective. Forty previously evaluated networks; retrospective sensitivity analysis.

Additional indirect-recording comparison

The held-out-prior versus nominal response-MSE change was -3.48% (95% CI -34.14% to +26.28%). The prespecified comparison did not establish an improvement from the held-out-prior criterion. The internal covariance criterion versus nominal gave -22.77% (95% CI -52.62% to -5.70%); prior-averaged local risk versus nominal gave -13.56% (95% CI -46.73% to -3.90%). All policies used the same budget and inference rule. The covariance-versus-ensemble comparison gave -10.65% (95% CI -18.73% to +3.93%); covariance versus random recording gave -5.50% (95% CI -24.37% to +1.72%). In the retrospective optimization sensitivity check, covariance versus nominal gave -11.50% (95% CI -17.17% to +15.04%). These tests concern an application of established covariance and predictive-validation ideas, not a new general experimental-design principle.

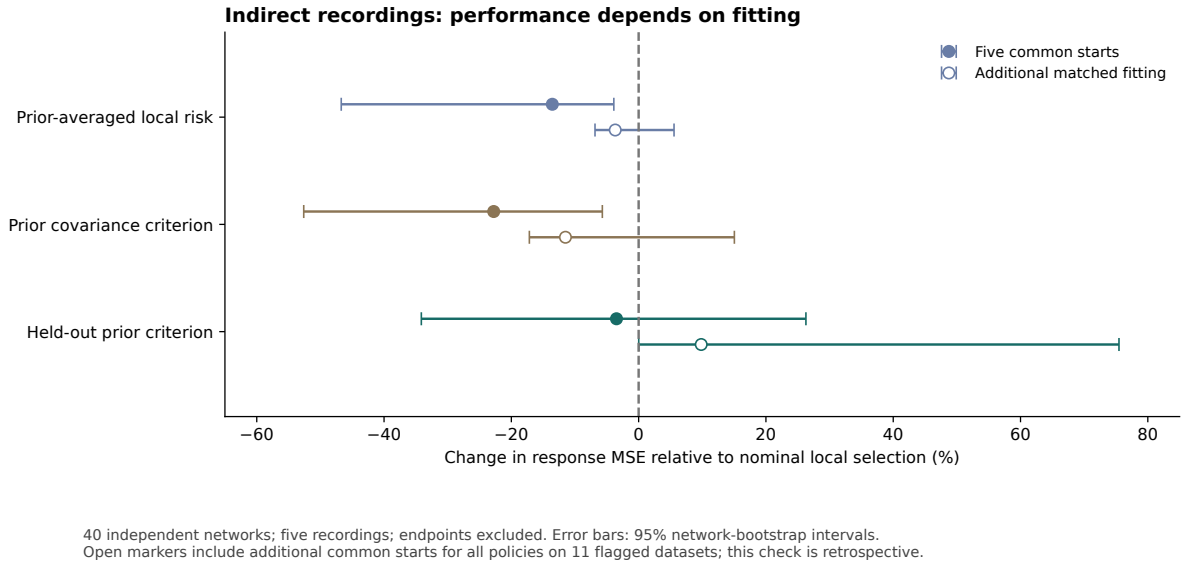


Fig. S7 | Indirect-recording criteria on forty independent networks. Points and intervals show paired mean-MSE changes and 95% network-bootstrap intervals relative to nominal local selection. Filled markers use five common starts; open markers include additional matched fitting for all policies on eleven flagged datasets. The original held-out-prior comparison is primary; other comparisons and the fitting sensitivity check are exploratory.